

Implementation of the unbiased KernelSHAP

SHAP

- **exact**
- computation time exponential in number of features (assuming blackbox model)

KernelSHAP

- approximates SHAP values via linear regression
- finds **exact** solution to the **approximate** problem by sampling
- consistent, $O(1/h)$ conv.
- *seems to be unbiased*
- empirically found convergence formula and error estimates

unbiased KernelSHAP

- approximates SHAP values via linear regression
- finds **approximate** solution to the **exact** problem by sampling
- consistent, $O(1/h)$ conv., slower than KernelSHAP
- provably unbiased
- higher variance than KernelSHAP
- derived convergence formula and error estimates

Italic font indicates news from the paper. Moreover a better sampling strategy is proposed and tested to lower estimators' variances.

[link to the paper](#)